



Systems Programming Concepts

Hic sunt dracones

MAJ John Rollinson

2023-03-08

Army Cyber Institute

Introductions

- Ask questions

- **Ask questions**
 - If you're confused now, so is someone else.

- **Ask questions**
 - If you're confused now, so is someone else.
 - If you're confused now, you will still be on the next slide.

- **Ask questions**
 - If you're confused now, so is someone else.
 - If you're confused now, you will still be on the next slide.
 - If you're confused now, the practical portions will hurt.

My Expectations

- **Ask questions**
 - If you're confused now, so is someone else.
 - If you're confused now, you will still be on the next slide.
 - If you're confused now, the practical portions will hurt.
- **Don't "cheat"**

My Expectations

- **Ask questions**
 - If you're confused now, so is someone else.
 - If you're confused now, you will still be on the next slide.
 - If you're confused now, the practical portions will hurt.
- **Don't "cheat"**
 - There's no grade

My Expectations

- **Ask questions**
 - If you're confused now, so is someone else.
 - If you're confused now, you will still be on the next slide.
 - If you're confused now, the practical portions will hurt.
- **Don't "cheat"**
 - There's no grade
 - Defeats the purpose of being here

My Expectations

- **Ask questions**
 - If you're confused now, so is someone else.
 - If you're confused now, you will still be on the next slide.
 - If you're confused now, the practical portions will hurt.
- **Don't "cheat"**
 - There's no grade
 - Defeats the purpose of being here
 - Won't work later

Course Overview

It's all ~~code~~ logic – not magic.

Abstractions leave gaps.

Own your tools.

- MAJ Chuck Suslowicz (possibly stolen from Roy Ragsdale)

*Manpages are your friend.*¹

- A previous C3T coach...

¹You're documenting your own code, right?

What to Expect

- Lots of information
- Lecture followed by “hands-on”
- Open-ended exercises (may not have “right” answers)

Course Objectives

- Understand the hardware/software interface
- Apply virtual memory concepts to userspace code
- Describe core OS concepts
- Examine code for concurrency issues
- Apply synchronization primitives & deadlock prevention techniques
- Understand distributed systems concepts


- Day 1: CPUs & Instructions
- Day 2: Virtual Memory
- Day 3: Operating Systems Overview
- Day 4: Concurrency
- Day 5: Distributed Systems Concepts

CPUs & Instructions


What *is* a computer?

What are the *components* of a computer?

Anatomy of a computer: Naive CPU



Anatomy of a computer: 2-Stage CPU



Anatomy of a computer: 5-Stage CPU

1. Fetch

Anatomy of a computer: 5-Stage CPU

1. Fetch
2. Decode

Anatomy of a computer: 5-Stage CPU

1. Fetch
2. Decode
3. Execute

Anatomy of a computer: 5-Stage CPU

1. Fetch
2. Decode
3. Execute
4. Memory

Anatomy of a computer: 5-Stage CPU

1. Fetch
2. Decode
3. Execute
4. Memory
5. Write-back

Anatomy of a computer: Stages Example

Initial state:

`rip` = 0x40000

`rsi` = 0x41000

`mem[0x41000]` = 0xDEADBEEF

1. **Fetch:** 48 8b 06 (@ address of *rip*)

Anatomy of a computer: Stages Example

Initial state:

`rip` = 0x40000

`rsi` = 0x41000

`mem[0x41000]` = 0xDEADBEEF

1. **Fetch:** 48 8b 06 (@ address of *rip*)
2. **Decode:** `mov rax, [rsi]`

Anatomy of a computer: Stages Example

Initial state:

```
rip          = 0x40000  
rsi          = 0x41000  
mem[0x41000] = 0xDEADBEEF
```

1. **Fetch:** 48 8b 06 (@ address of *rip*)
2. **Decode:** mov rax, [rsi]
3. **Execute:** mov rax, [0x41000]

Anatomy of a computer: Stages Example

Initial state:

```
rip          = 0x40000  
rsi          = 0x41000  
mem[0x41000] = 0xDEADBEEF
```

1. **Fetch:** 48 8b 06 (@ address of *rip*)
2. **Decode:** `mov rax, [rsi]`
3. **Execute:** `mov rax, [0x41000]`
4. **Memory:** `mov rax, 0xDEADBEEF`

Anatomy of a computer: Stages Example

Initial state:

```
rip          = 0x40000  
rsi          = 0x41000  
mem[0x41000] = 0xDEADBEEF
```

1. **Fetch:** 48 8b 06 (@ address of *rip*)
2. **Decode:** `mov rax, [rsi]`
3. **Execute:** `mov rax, [0x41000]`
4. **Memory:** `mov rax, 0xDEADBEEF`
5. **Writeback:** `rax = 0xDEADBEEF`

Anatomy of a computer: Pipelining Example

CPU

Fetch
Decode
Execute
Memory
Writeback

Data (0x1000)

1
42
1337
0

Instructions

00050593
00052503
02050063
00000513
00458593
0005a603
00150513
00458593
fe061ae3
00008067
00000513
00008067

Registers

x0	0	1	2	3	4	5	6	7
x8	8	9	0x1000	11	12	13	14	15
x16	16	17	18	19	20	21	22	23
x24	24	25	26	27	28	29	30	31

Anatomy of a computer: Pipelining Example

CPU

Fetch	00050593
Decode	
Execute	
Memory	
Writeback	

Registers

x0	0	1	2	3	4	5	6	7
x8	8	9	0x1000	11	12	13	14	15
x16	16	17	18	19	20	21	22	23
x24	24	25	26	27	28	29	30	31

Data (0x1000)

1
42
1337
0

Instructions

00050593
00052503
02050063
00000513
00458593
0005a603
00150513
00458593
fe061ae3
00008067
00000513
00008067

Anatomy of a computer: Pipelining Example

CPU

Fetch	00052503
Decode	mv a1, a0
Execute	
Memory	
Writeback	

Registers

x0	0	1	2	3	4	5	6	7
x8	8	9	0x1000	11	12	13	14	15
x16	16	17	18	19	20	21	22	23
x24	24	25	26	27	28	29	30	31

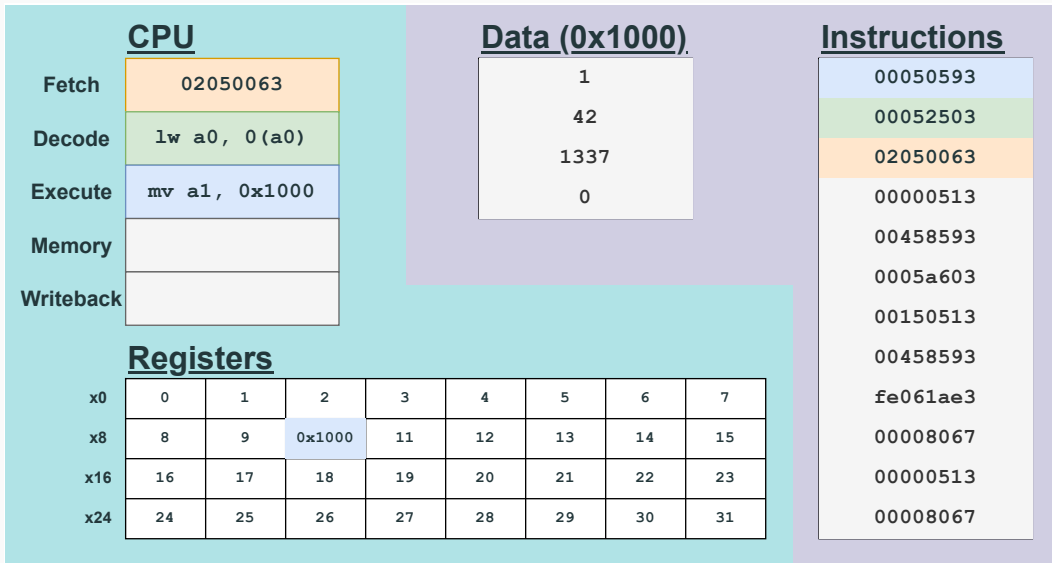
Data (0x1000)

1
42
1337
0

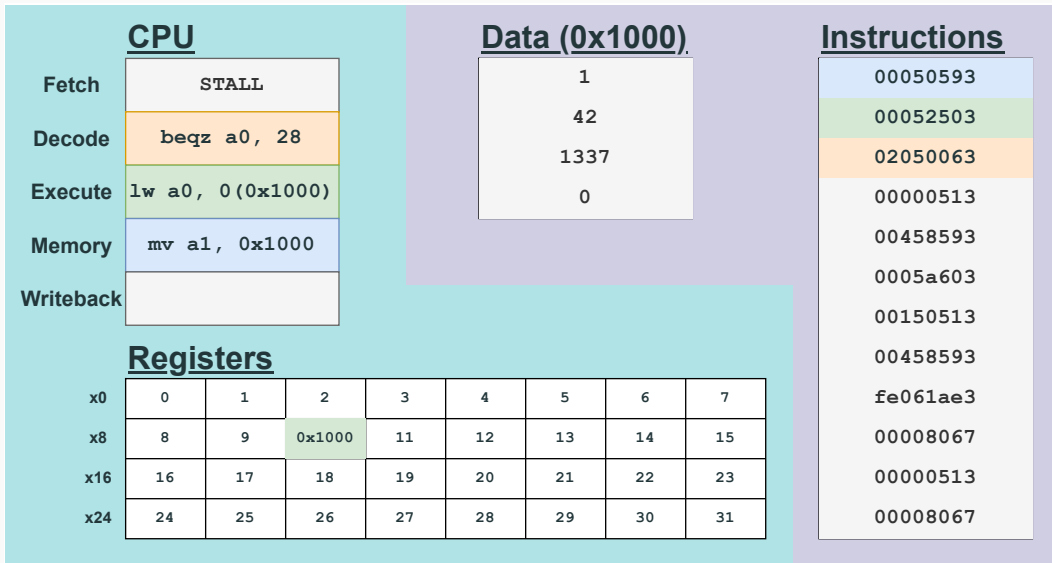
Instructions

00050593
00052503
02050063
00000513
00458593
0005a603
00150513
00458593
fe061ae3
00008067
00000513
00008067

Anatomy of a computer: Pipelining Example



Anatomy of a computer: Pipelining Example



Anatomy of a computer: Pipelining Example

CPU

Fetch	STALL
Decode	beqz a0, 28
Execute	STALL
Memory	lw a0, 1
Writeback	mv a1, 0x1000

Registers

x0	0	1	2	3	4	5	6	7
x8	8	9	0x1000	0x1000	12	13	14	15
x16	16	17	18	19	20	21	22	23
x24	24	25	26	27	28	29	30	31

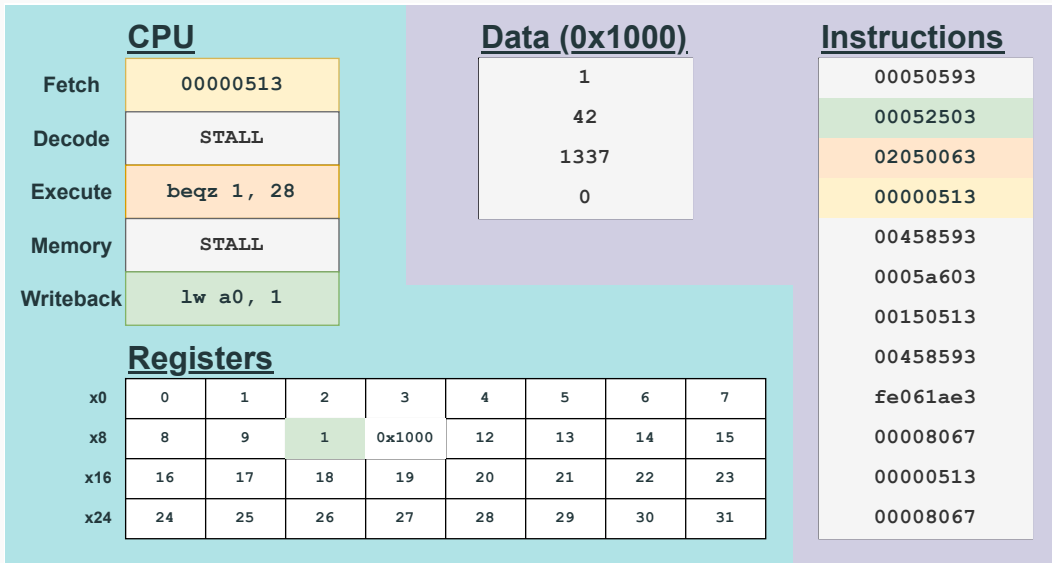
Data (0x1000)

1
42
1337
0

Instructions

00050593
00052503
02050063
00000513
00458593
0005a603
00150513
00458593
fe061ae3
00008067
00000513
00008067

Anatomy of a computer: Pipelining Example



Anatomy of a computer: Check on Learning

```
00: 00050593      mv    a1, a0
04: 00052503      lw     a0, 0(a0)
08: 02050063      beqz  a0, 16 <.LBB0_4>
0c: 00000513      li     a0, 0
10: 00458593      addi   a1, a1, 4
.LBB0_2:
14: 0005a603      lw     a2, 0(a1)
18: 00150513      addi   a0, a0, 1
1c: 00458593      addi   a1, a1, 4
20: fe061ae3      bnez  a2, -6 <.LBB0_2>
24: 00008067      ret
.LBB0_4:
28: 00000513      li     a0, 0
2c: 00008067      ret
```

Instruction Set Architecture (ISA): The abstract model of a computer. It encompasses the instructions, registers, memory, and I/O interfaces.

Machine Code: The machine readable representation of executable instructions.

Assembly Language: A human-friendly representation of machine code.

Reduced Instruction Set Computers (RISC): each instruction performs a single operation

Complex Instruction Set Computers (CISC): some instructions combine two or more operations to improve code density

Note: Reality is more complicated than this.

ISAs may not target “real” hardware and may target a software-defined runtime (e.g., JVM or eBPF).

The compiled code is often referred to as *bytecode*.

Sometimes called “managed” code (i.e., CLR on Windows).

Instruction Sets: Writing Assembly

inline.c

```
bool checked_add(int64_t *result,
                 int64_t  left,
                 int64_t  right) {
    int overflow = 0;
    asm volatile (
        "add %0, %2;\n"
        "adox %1, %1;\n"
        : "+r" (left), "+r" (overflow)
        : "r" (right));
    *result = left;
    return overflow;
}
```

separate.S

```
.global checked_add
checked_add:
    xor rax, rax
    add rsi, rdx
    adox eax, eax
    mov [rdi], rsi
    ret
```

- [Intel Software Developer Manuals](#)
- [AMD64 Architecture Manual](#)
- [ARM Documentation](#)
- [RISC-V Specification](#)
- [JVM Specifications](#)

The AMD64 ISA will be the focus of the remainder of this class when specifics are needed. Other ISAs differ in important details but most of the principles are transferable.

Registers: `mov rbp, rsp`

Immediates: `mov rax, ffh`

Memory:

- `push rbp`
- `pop rbp`
- `mov [rdi], rdx`
- `mov rax, [rbp - 8]`
- `mov rax, [rdi + rsi*8]`
- `mov rax, [rdi + rsi*8 + 24]`

Generic form: `OP DST, SRC` $\implies dst = op(dst, src)$

Special considerations:

- lock prefix required for atomic update of a memory address
- `lea rax, [rdi + rsi*8 + 24]` $\implies rax = rdi + 8 * rsi + 24$
- *EFLAGS* register has special flags that get updated.

Relative control flow changes:

- `jmp 8`
- `call 88h`
- `jnz -8`

Absolute control flow changes:

- `jmp rax`
- `call rax`
- `call [rdx]`
- `ret`

Arrays:

```
int array[8];  
array[x] = 42;  
// mov [rdi + rsi*4], 42  
//      array + x * sizeof(int)
```

Structs:

```
struct {  
    uint16_t a;  
    uint8_t b;  
    uint8_t c;  
} s;  
s.c = 42;  
// mov [rdi + 3], 42  
//      &s + offsetof(s, c)
```


Exercise 1.1

Reading/writing assembly (demo).

Anatomy of a function

- Prologue
- Epilogue
- Calling conventions
 - Args
 - Return
 - Caller-saved
 - Callee-saved
 - *Alignment*

Anatomy of a function: Prologue

```
; Create the stack frame (present most of the time)
push rbp
mov  rbp, rsp
sub  rsp, 64
; Save Callee-saved variables
mov  [rbp - 40], rbx
mov  [rbp - 48], r12
; Debugging support
mov  [rbp - 8], rdi
mov  [rbp - 16], rsi
mov  [rbp - 24], rdx
mov  [rbp - 32], rcx
```

Anatomy of a function: Epilogue

```
mov r12, [rbp - 48]
mov rbx, [rbp - 40]
leave
; leave is equivalent to `mov rsp, rbp; pop rbp`
ret
```

Anatomy of a function: Calling Convention: Windows

Arguments:

- rcx
- rdx
- r8
- r9
- 32-bytes reserved stack
- Stack “right-to-left”

Return:

- rax

Volatile:

- All arguments
- rax
- r10
- r11

Anatomy of a function: Calling Convention: AMD64

Arguments:

- rdi
- rsi
- rdx
- rcx
- r8
- r9
- Stack “right-to-left”

Return:

- rax
- rdx

Volatile:

- All arguments
- rax
- r10
- r11

Exercise 1.2

1. Write `fibonacci` in recursive assembly ($f(n) = f(n - 1) + f(n - 2)$).
2. Write a function that prints the address after the *call-site* (calls `printf`).
3. Write a function that prints the value of all registers at the call-site.
4. Write a function that prints as many return addresses in the call-stack as it can. (It should not segfault)

See `checker12.c` in `ex12.tar.gz` for a simple test harness.

Traps, Interrupts, & Exceptions

Trap: deliberate transition into the kernel

Interrupt: external events (timers, I/O, etc.)

Exception: error occurred

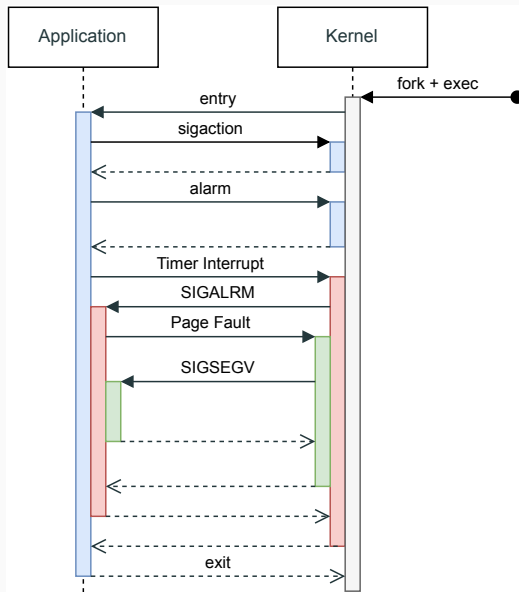
Signal: kernel concept

Traps, Interrupts, & Exceptions: Alarm: Code

```
void alarm_handler(int unused) {  
    int *p_bad = (int *)0x1;  
    *p_bad = 42;  
}
```

```
void segfault_handler(int unused, siginfo_t *info, ucontext_t *uctx) {  
    /* Platform-specific code to change address of `p_bad` */  
}
```

Traps, Interrupts, & Exceptions: Alarm: Diagram



Exercise 1.3

Writing syscalls directly:

1. Give code that opens “/etc/passwd”, reads it in, and writes it to STDOUT
2. Exec a program to print the flag
3. Fork and execute a program in the background (e.g. socat).

Privilege Levels

- Rings / kernel mode
- Why are they necessary/helpful?
- What do they protect? (I/O & memory)

Speaking of memory. . .