



# Systems Programming Concepts

*Hic sunt dracones*

---

MAJ John Rollinson


2023-03-08

Army Cyber Institute

# Virtual Memory

---


# Anatomy of a computer



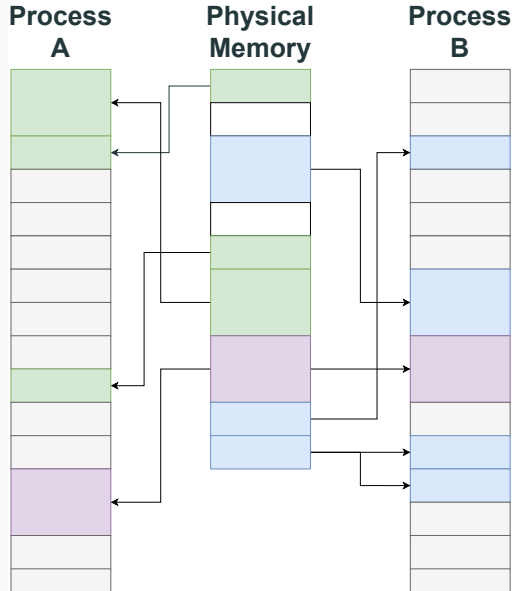
# Anatomy of a computer: Managing Memory

What's the problem with this picture?

# The Memory Management Unit



# The Memory Management Unit: Logical View



# The Memory Management Unit: Implementations

- Segments<sup>1</sup>
- Page tables
  - Hashtable
  - Trees
- Software defined

---

<sup>1</sup>This isn't necessarily the same, but can achieve similar results.

# The Memory Management Unit: Segments

- Effectively multiple address spaces
- Segment is a base address and length in memory
- Segments identified by extra machine registers

$$addr_p = addr_v + base_{segment}$$



## The Memory Management Unit: Hashtables

1. Lookup  $addr_v$  in a small table of segments.
2. (Optional) Check a hashtable for the segment info.
3. Fault the CPU and have the OS update the segment tables.


*Note:* Only really used by POWER chips.

# The Memory Management Unit: Trees

- Lowest-order bits are the offset into a page.
- Higher-order bits are used to walk a radix tree.

# The Memory Management Unit: x86-64 Details

## Virtual Address



## Exercise 2.1

Exercise: Write a library that given a “physical memory dump” and CR3 value, and whether the CPU is in supervisor-mode, can dereference logical values and update memory.

See `mmu.h` in `ex21.tar.gz`.

How many physical memory accesses does this instruction require?

```
call [rdi]
```

Fetching memory from main memory takes hundreds of CPU cycles on modern CPUs. To reduce this effect, they often have **caches** that exploit **locality** of memory accesses to retain data “closer” to the CPU where accesses can be as fast as 2-5 cycles.

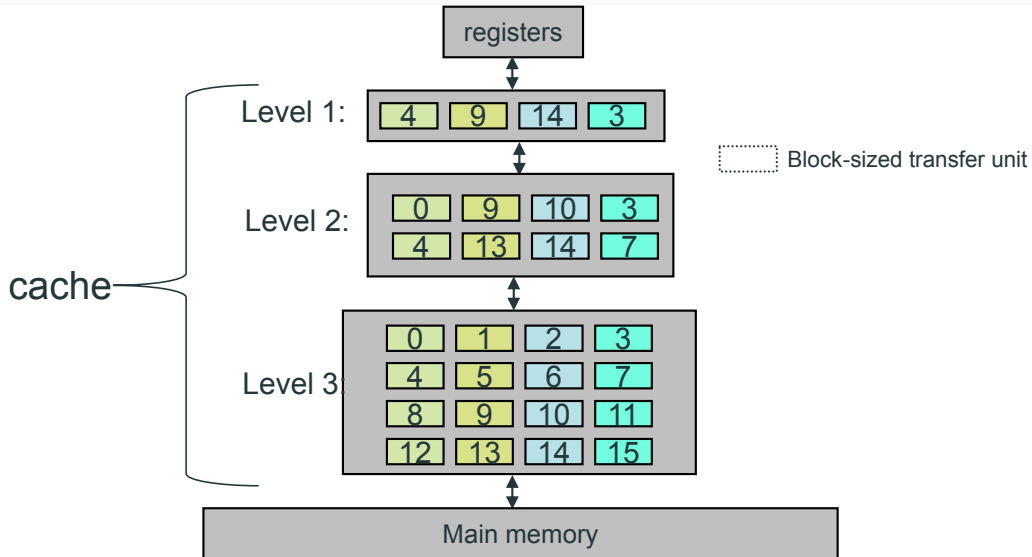
## Caches: Locality

***Spatial locality:*** Addresses of memory accesses are close

***Temporal locality:*** The same address is accessed repeatedly


```
size_t checksum(const uint8_t *data, size_t size) {  
    size_t checksum = 0;  
    for (size_t index = 0; index < size, index++) {  
        checksum *= SOME_CONST;  
        checksum += data[index];  
    }  
    return checksum;  
}
```

# Caches: Model





# Caches: Anatomy of a computer



## Caches: Computational Complexity

Should we count loads/stores, comparisons, or something else? Why?

When the cache is full, does our replacement strategy matter?

## Translation Lookaside Buffers (TLBs)

- Does the cache fully solve our memory translation problems?


## Translation Lookaside Buffers (TLBs)

- Does the cache fully solve our memory translation problems?
- TLBs are special caches for page mapping

## Translation Lookaside Buffers (TLBs)

- Does the cache fully solve our memory translation problems?
- TLBs are special caches for page mapping
- Separate paths for instructions and data

# Translation Lookaside Buffers (TLBs): Anatomy of a computer



Why does all of this matter?

## Exercise 2.2

Write programs to measure and determine (or approximate) the following information:

1. A cache line
2. The L1 cache (Bonus points for L1 iCache)
3. The L2 cache
4. The L3 cache
5. TLB size (bonus points for iTLB size)



# Storage hierarchy

